

Network Simulation with **ns/2**: report

Maarten Tegelaers Thomas Vochten

May 1, 2014

Abstract

This report describes the use of the **ns/2** network event simulator to simulate the behaviour of simple network topologies. Two topology situations are considered. For the first situation, the focus lies on an **FTP!** (**FTP!**) application competing for bandwidth with a **CBR!** (**CBR!**) application. In the second situation, we investigate the evolution of the congestion window size during the execution of the **TCP!** (**TCP!**) algorithm in detail; an **FTP!** application is running continuously while a web user on the same LAN as the **FTP!** application generates short intense bursts of traffic. We also compare congestion window size behaviour between the Tahoe implementation of **TCP!** and the Reno implementation of **TCP!**.

1 Exercise 1: Bandwidth restrictions on KotNet

1. First, leave out the connection of the uploader. Plot the throughput of the main **FTP!** connection and discuss the download behaviour of a KotNet/Telenet modem with bandwidth caps. How do bandwidth limitations affect the download connection?
2. What happens if you now enable both upload and download activities? Plot and compare the throughput of both streams. Do you spot any problems? If so, explain.
3. Consider now that the model is able to guarantee a certain amount of upload bandwidth to every application, which performance do you expect for both applications (**FTP!** and **CBR!**)?
4. What would be the result if bandwidth was even more restricted? For example, up to September 2010, KotNet limitations were 512Kbps for downstream traffic and 100Kbps for upstream traffic.
5. Configure the upload connection with a fixed `rate_` of 30k. Explain the simulated result (no plot required). What would be your (extra) solution(s) to improve the download experience on shared LANs behind capped Internet connections? How could you practically realise this?
6. Suppose we have in the future a KotNet/Telenet cable subscription that has 10 times more capacity: i.e a connection that has a downstream capacity of 40Mbps with an upstream capacity of 2Mbps.

Figure 1: Throughput

- (a) What performance can you expect if there are 10 users behind this connection and all performing the same activities as in this exercise? Why? What performance can you expect if there are only 5 users using the network? Why?
- (b) Does the performance differ if 10 users perform the same activities but at random times? Are there some users who suffer from the network activities caused by other users? Why?

2 Exercise 2: Tahoe and Reno versus bursty web traffic

1. **Investigate the throughput of the main FTP application. Can you indicate the effects the individual bursts of web traffic have on its total throughput? Why do these effects not immediately become visible when each burst starts (i.e. at 5, 10, or resp. 15 seconds)?**

Figure ?? shows the throughput of the main FTP application. We can see that the throughput rises rapidly at the start before hitting the bandwidth limit within the first 700 milliseconds. Thereafter, it remains at 1.25 megabytes/s. When the first burst begins at 5 seconds, we see a drop to 800 kilobytes/s as the first TCP connection for the web user starts, but the throughput can still climb to one megabytes/s. It is only after 700 milliseconds that the throughput decreases dramatically. When the burst starts, new connections between the web user and the web server are opened within about 0.05 milliseconds. These connections first go through the slow start phase of TCP, which means it takes some time before each connection attains a throughput that could affect existing connections. Soon, the 10 megabytes/s bandwidth limit cannot support high throughputs for each active connection, which means the main FTP connection experiences packet loss. This triggers congestion avoidance for the main FTP application and it sets its congestion window to 1 packet, after which it must go through slow start again.

After about 8 seconds, the main FTP application has to contend with increasingly fewer connections from the web user, which allows the throughput to climb steadily. A new burst starts at 10 seconds, and this time the effects manifest themselves after 300 milliseconds. This is again explained by connections being opened one after another and each connection gradually increasing its throughput.

Right before the third burst, the FTP application cuts its throughput because of congestion. The throughput is allowed to climb until 15.7 seconds, when the new burst first has an impact. From 17 seconds on, more

and more web connections finish their work, allowing the FTP application to steadily increase its throughput to the highest allowed bandwidth across the network.




Figure 2: The throughput of the continuously running FTP application. Ticks are drawn every 200 milliseconds.

2. **Plot the congestion window size and slow start thresholds in another graph. Can you identify the different phases of the TCP algorithm? When is the slow start threshold recalculated and how?**

Figure ?? shows the congestion window sizes and the slow start thresholds for the main FTP application. To start off, the threshold is set to 80 packets and the congestion window size to 1 packet. Each time the sender receives an acknowledgement, the congestion window is set to $\min(2 \cdot \text{window size}, \text{threshold})$. This is the slow start phase. Once the threshold has been reached, the window size increases by one each round-trip time. This is the additive increase phase. At about 5.76 seconds, the FTP application experiences packet loss, which causes the window size to drop to one packet. The slow start threshold is set to $\max(2, \frac{\text{threshold}}{2})$, which in this case is 40. After that, the slow start phase begins again.

3. **Discuss the AIMD principle of TCP. Take one 'sawtooth' pattern in the graph and indicate a few reasonable values for the congestion window on the graph. What is the first interval the TCP congestion avoidance algorithm is active?**

AIMD stands for Additive Increase Multiplicative Decrease. It is the idea




Figure 3: The congestion window sizes and slow start thresholds for the main FTP application running TCP Tahoe.

where one wants to increase the congestion window size by one each round-trip time and, when necessary, decrease the window size multiplicatively. In TCP Tahoe, this multiplicative decrease is implemented by cutting the window size to one packet if packet loss occurs. TCP diverges from the principle by the use of slow start, where the window size increases exponentially (by always doubling the window size) until a pre-defined threshold is reached. After that, it goes back to additive increase. In TCP Tahoe, slow start is executed often, while in TCP Reno this only happens when a connection is first established and whenever fast recovery fails, which is explained later on in this section.

When one plots the behaviour exhibited by a connection that abides by the AIMD principle, it produces a sawtooth pattern. There are periods where the congestion window size increases, overshoots the optimal window size, decreases, falls below the optimal window size and then increases again. The effect of applying AIMD is that the congestion window size converges to the optimal value. Looking at the moment where the first burst begins at 5 seconds in ??, there is a first drop to a size of one packet. After that, slow start is activated again and the window size increases rapidly to 40 packets, after which congestion immediately manifests itself again. The next slow start increases until a size of 20, after which additive increase continues for about 600 milliseconds when there is packet loss again. After this drop, the intensity of the current burst decreases. This behaviour suggests that a congestion window size of 20 would have been a reasonable value during this burst.

For the first burst, the congestion avoidance algorithm is active from about 5.76 seconds to 7.95 seconds, whereafter additive increase continues unim-

peded until after the start of the next burst.

4. **Change the TCP implementation of the main FTP application into Reno. Plot the congestion window and slow start thresholds. Do you spot the differences with the default TCP Tahoe implementation? Why does the window sometimes still drop to zero or one?**

Figure ?? shows the behaviour for the TCP Reno implementation. Looking at the congestion drop that occurs at approximately 10.6 seconds, it can be seen that both the congestion window size and the slow start threshold are set to $\max(2, \frac{threshold}{2})$. After that, the window size increases additively. Reno differs in this from Tahoe, because Tahoe drops the congestion window size to one packet and then runs a slow start until the new threshold is reached.

Both Tahoe and Reno take action in order to prevent congestion as well as possible if three duplicate acknowledgements are received for the same packet. Reno halves the congestion window. After that it enters fast recovery mode, where it resends the lost packet and waits for an acknowledgement. If the acknowledgement for the resent segment times out, the congestion window size is set to one segment and slow start is activated. This explains why the congestion window size is still sometimes set to one.




Figure 4: The congestion window sizes and slow start thresholds for the main FTP application running TCP Reno.